

<https://www.wsj.com/tech/ai/ai-computing-power-demand-a63da9b9>

TECHNOLOGY • ARTIFICIAL INTELLIGENCE

The Desperate Hunt for AI Computing Power Is Upending Silicon Valley

Bitter rivals are forming alliances and executives are having to personally intervene over the scramble for resources to fuel the AI boom

By Kate Clark, Erin Woo, Keach Hagey and Meghan Bobrowsky

Oct. 9, 2026 at 9:00 pm ET



TERMS & CONDITIONS FOR WSJ, GETTY IMAGES

Google really wanted pioneering AI researcher Noam Shazeer back.

The longtime employee left to start his own company but in 2024, Google paid

\$2.7 billion for him to return to work on cutting-edge projects, including its most powerful AI technology, Gemini.

Even that, though, couldn't ensure that Shazeer got the computing resources he wanted.

Less than two years later, he sent a note to one of his teams saying that Google had cut what was allocated to them in favor of another project. Shortly thereafter, Shazeer left for rival OpenAI.

Compute, the industry shorthand for computing power, is the lifeblood of the AI boom, and it's in short supply. Now fights over how to get more of it—and what projects to give priority to—are sending a powerful industry scrambling. Bitter rivals are forming new alliances, such as Elon Musk and Dario Amodei, whose companies worked together to secure computing power. Venture-capital firms are touting access to AI chips to woo startups. And power struggles are more frequent as companies compete in a worldwide treasure hunt for the tech industry's most coveted resource.

“There's a big gold rush happening at the moment,” said Evan Conrad, chief executive officer of San Francisco Compute. “It's hard to fully comprehend the scale.”

Conrad's company vets data centers and then rents their computing power to AI companies that range from the largest players like Nvidia to small startups like Standard Intelligence. Its chief technology officer, Eric Park, travels the globe in search of computing resources hiding where few else think to look.

In the spring, it was far-flung corners of rural America, where one promising new data center turned out to be AI servers humming inside what appeared to be a converted chicken coop. Then Manila in monsoon season, where Park's car stalled in floodwater and he and his colleagues waded to safety, bags held over their heads. Eventually, a passing truck shuttled him to the airport, where he

bought a new pair of shoes and boarded a flight to Taipei—to meet with a company building servers for the AI boom.



Eric Park traveling on foot through a monsoon in Manila in search of computing resources. ERIC PARK

Even a little cluster of available AI chips can be helpful to startups trying to train or fine-tune smaller AI models.

SF Compute's business has grown as companies are increasingly beholden to the yearslong timelines of data-center build-outs. There's a race to build more data centers to get more capacity online fast, but construction is stymied by regulation and political backlash.

Established suppliers of computing power are requiring customers to sign multiyear contracts with hefty down payments. Startups are increasingly inking deals with unreliable providers out of desperation, like ones lacking proper security or weaker infrastructure.

While splashy announcements have promised enormous amounts of new computing power coming soon, far less is actually available today, said Nick Frosst, co-founder of Canadian AI startup Cohere.

“Lots of things have been announced and then canceled,” said Frosst, whose startup builds large language models and was valued at \$6.8 billion last year. “Or things that they said they’re going to build this big, but then they actually build significantly less.”

Arcee, an open-source AI startup, recently secured commitments for both capital and computing capacity from an infrastructure partner, said CEO Mark McQuade, but the deal collapsed when the firm’s data center was delayed.

“So much is up in the air with these things, especially when it comes to compute,” said McQuade. “You take it with a grain of salt until it’s real.”

With so much demand, prices are going up. The hourly cost of renting computing power from Nvidia H100 chips on a one-year contract has risen 60% in the past year, according to SemiAnalysis, a data and consulting firm. Alphabet **GOOGL 0.97%** ▲ expects to make up to \$200 billion this year in capital expenditures, much of it going toward AI infrastructure.

Even the busiest AI executives are personally involving themselves in negotiations over computing power.

In late March, Anthropic co-founder Tom Brown visited Elon Musk at rival AI company xAI’s offices in Hawthorne, Calif., to broker a deal to rent computing power from Musk’s company, according to people familiar with the matter. The explosive, unexpected success of Anthropic’s Claude Code had created a computing crunch, contributing to frequent outages.

In May, Anthropic stunned the industry when it announced an agreement with Musk’s SpaceX to lease more than 300 megawatts of computing power—part of a deal later revealed to be worth up to \$45 billion in the coming years. Musk, who had previously been publicly critical of Anthropic, posted on X that he was “impressed” with members of Anthropic’s senior team, adding, “No one set off my evil detector.”

Anthropic CEO Dario Amodei also spoke on the phone with Meta chief AI officer Alexandr Wang earlier this year, attempting to source more computing capacity, according to a person familiar with the matter. Meta debated whether to give chips to Anthropic and ultimately decided not to for now, people familiar with the matter said.

This is all in addition to the hundreds of billions of dollars in computing deals Anthropic has announced over the past year.



Dario Amodei KEVIN DIETSCH/GETTY IMAGES

Inside Google's Mountain View offices, researchers are fiercely divided over

which projects deserve the company's precious computing resources. Co-founder Sergey Brin sometimes upends the formal compute allocation process, reshuffling power around to projects that he thinks are important, according to people familiar with the matter. Demis Hassabis, who co-founded and until recently ran Google's AI lab DeepMind and took a chairman role in August, has expressed frustration about a lack of compute allocated toward long-term research, according to people familiar with the matter.

To hedge against the unpredictable nature of securing computing power, some companies are signing contracts for more capacity than they need. In other cases, startups have predicted surging demand for their AI products that never materialized, forcing them to resell leftover capacity.

This confluence of factors coupled with rising wait times and skyrocketing costs of compute are crippling a generation of newcomers trying to compete with Google, Anthropic and OpenAI.

Earlier in the AI boom, companies large and small could rent clusters of AI chips from Amazon Web Services, for example, without having to make a multiyear commitment or pay an astronomical upfront fee.

As frontier AI labs soaked up the available compute, cloud providers moved to require multiyear contracts with upfront costs often as high as 30%. Sometimes they require startups to find guarantors willing to cover their payments if they are unable to.

"It's like if you had an apartment, and you could rent month-to-month, and then suddenly your landlord comes in and says, 'Hey, I need you to sign a five-year lease, and by the way, I need you to pay a year upfront,'" said SF Compute's Conrad.

The challenge isn't only securing computing power but accurately predicting how much firepower startups will need as they scale.

One of Silicon Valley's newest and fastest-growing startups, the viral AI assistant Instinct, is dealing with steep competition from Meta and OpenAI, which both recently debuted their own personal AI agents and have access to substantially more capital and computing power.

Instinct last month raised \$1 billion at a \$10 billion valuation in what was its third capital raise this year, but even with oodles of cash, it may struggle to keep pace with the likes of Meta, which owns and operates its own data centers.

Instinct CEO Noah Shinn said on a recent podcast episode of "Invest Like the Best" that he has wrestled with whether to buy compute now to accommodate a future where it has 100 million users, a risky move that could tank the company if it doesn't grow as expected. A representative for Instinct declined to comment.

"If you overbuy, you could bankrupt the company by having too many lease commitments, and if you underbuy, you could knock down your growth trajectory and have a bunch of unhappy users," said Zach Bratun-Glennon, a general partner at Gradient Ventures, which invests in AI startups.

Some venture-capital firms are now reserving clusters themselves and renting the capacity back to their startups, or offering to guarantee contracts with compute providers. Last year, one of Y Combinator's startups needed 120 graphics processing units, or GPUs, for a two-month training run but couldn't find a suitable contract, according to a person familiar with the accelerator. So Y Combinator introduced a cluster dedicated to its startups, which it launched in July.

At Radical Ventures, an AI-focused VC firm that recently backed a new AI lab co-founded by Alphabet's former chief scientist Jeff Dean called Discovery Loop, David Katz leads a team that tracks available GPUs, monitors pricing and contract terms, and connects portfolio companies with compute providers.

David Katz RADICAL VENTURES

Over the past five years, Katz and his small team have negotiated billions in compute deals, he said. Katz often finds himself texting directly with the leaders of neoclouds—companies that purchase GPUs and other chips then lease them to customers—to inquire about availability. He flags potential investments to compute partners before Radical has decided to invest to make sure capacity will be there if they need it.

“It used to be, ‘Where am I getting my capital?’ Now it’s like, ‘Who’s going to help me?’” said Katz. “There’s a lot of capital floating around for these companies, but there’s not a lot of deep intelligence on how to build these things.”

Kate Clark covers startups, venture capital and artificial intelligence for The Wall Street Journal and is based in New York. Her reporting examines venture investment, private market dealmaking and the power dynamics between founders and investors in Silicon Valley and beyond. Previously, Kate was a senior...

Follow



Erin Woo is a reporter in The Wall Street Journal's San Francisco bureau, where she covers Google. She writes about how Google is reshaping itself and its products in the AI era, and how that reverberates across our online and offline lives. Erin previously covered Google, AI, social media and venture capital at The...

Follow



Keach Hagey is a reporter at The Wall Street Journal covering the intersection of media, technology and power. Her reporting explores how institutions and individuals wield influence in the new information economy, with a current focus on artificial intelligence and OpenAI. She is the author of "The Optimist: Sam...

Follow



Meghan Bobrowsky is a reporter in The Wall Street Journal's San Francisco bureau, where she focuses on Meta Platforms, artificial intelligence and social media. She writes about the company's massive data center buildout, its social media platforms that have billions of users, and the AI products the company is...

Follow



Videos

